

# TrustyAI x FMS Guardrails x Llama Stack

# 👋 About me – Mac Misiura

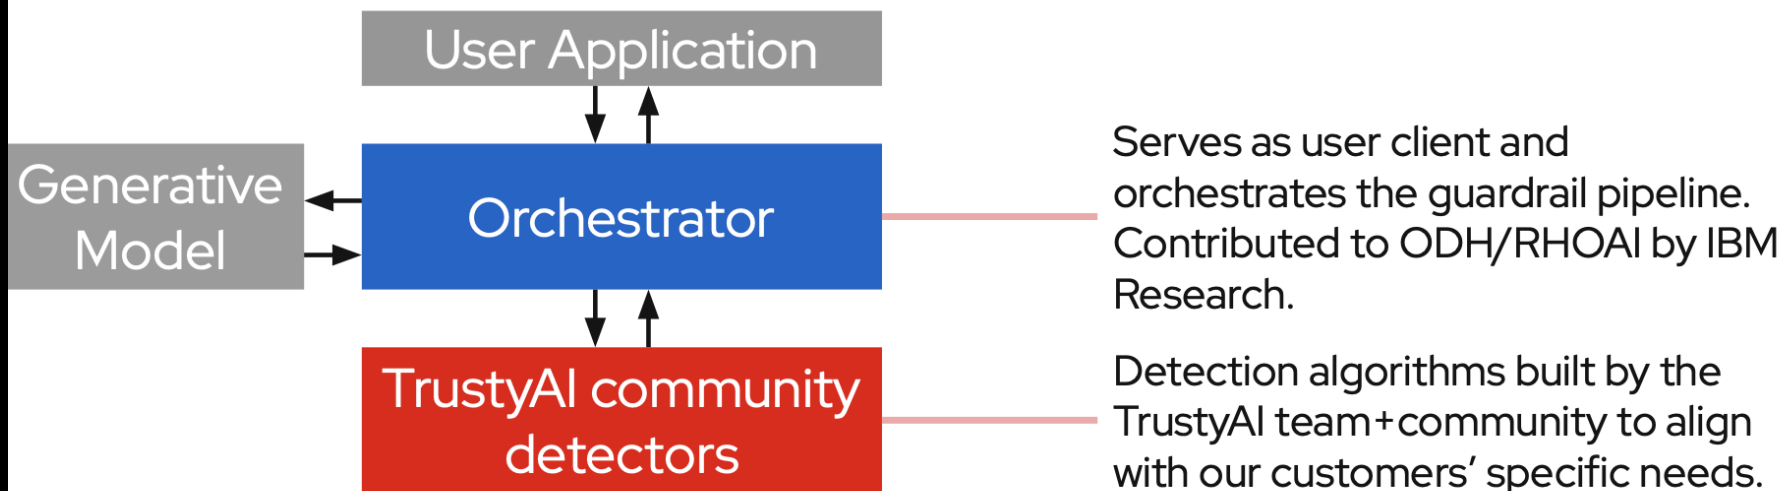


- obtained a PhD in Applied Mathematics and Statistics from Newcastle University in 2021
- previously worked as a Data Scientist (NLP) at [the National Innovation Centre for Data](#)
- joined [TrustyAI](#) as a Software Engineer (Machine Learning) in August 2024
- currently working on the Guardrails project and its integration with Llama Stack



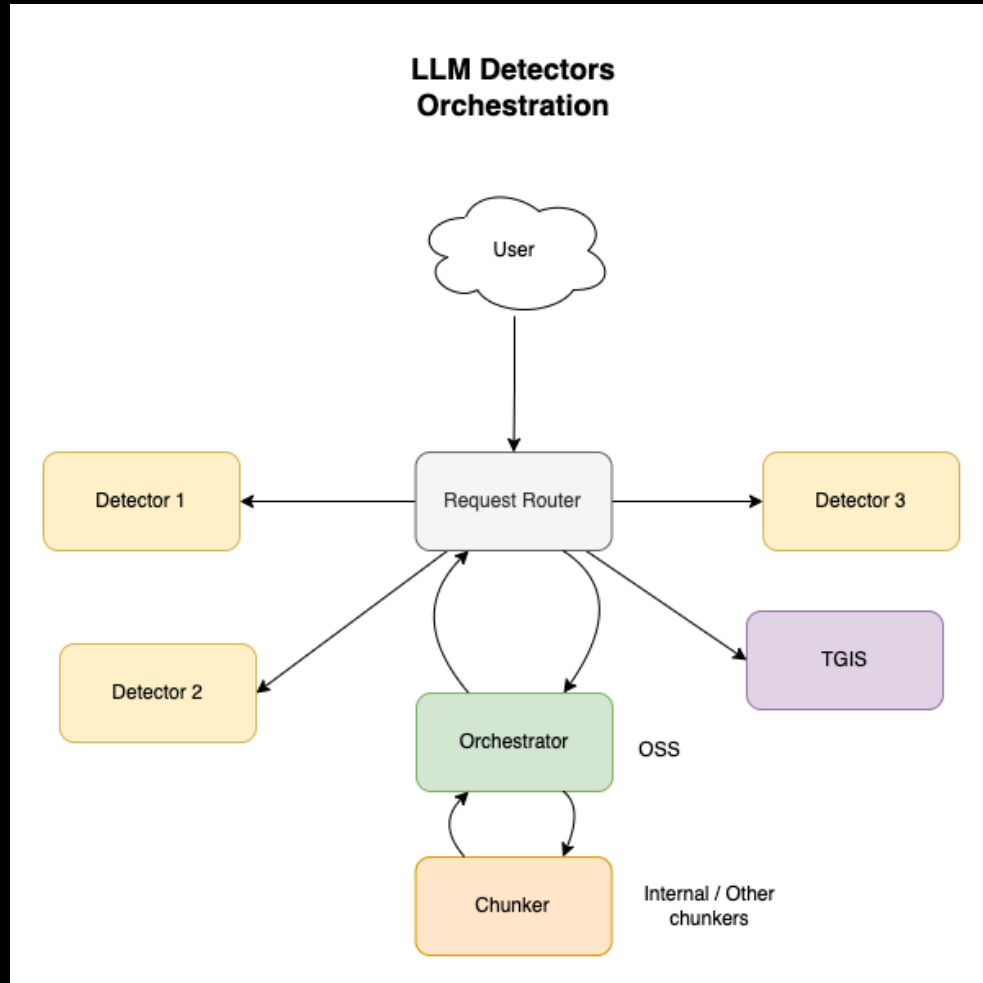
# FMS Guardrails – (sky) high level architecture

## Architecture





# Core component: the orchestrator



- the orchestrator has been implemented as a component of the [TrustyAI Kubernetes Operator](#)
- for information on how to get started, check out this [doc](#)

Image credit: [fms-guardrails-orchestrator repo](#)



# **Core component: the community detectors**

At present, the following community detectors are available:



# Orchestrator API



Select a definition

Orchestrator API

# FMS Orchestrator API

0.1.0

OAS 3.0

[docs/api/orchestrator\\_openapi\\_0\\_1\\_0.yaml](docs/api/orchestrator_openapi_0_1_0.yaml)

## Task - Text Generation, with detection

Detections on text generation model input and/or output

POST	<b>/api/v1/task/classification-with-text-generation</b>	Guardrails Unary Handler	▼
POST	<b>/api/v1/task/server-streaming-classification-with-text-generation</b>	Guardrails Server Stream Handler	▼
POST	<b>/api/v2/text/generation-detection</b>	Generation task performing detection on prompt and generated text	▼

## Task - Detection

Standalone detections

POST	<b>/api/v2/text/detection/content</b>	Detection task on input content	▼
POST	<b>/api/v2/text/detection/stream-content</b>	Detection task on input content stream	▼
POST	<b>/api/v2/text/detection/chat</b>	Detection task on entire history of chat messages	▼



# Detectors API





Select a definition

Detector API

# Detectors API

0.0.1

OAS 3.0

docs/api/openapi\_detector\_api.yaml

Apache 2.0

## Text Detections on text

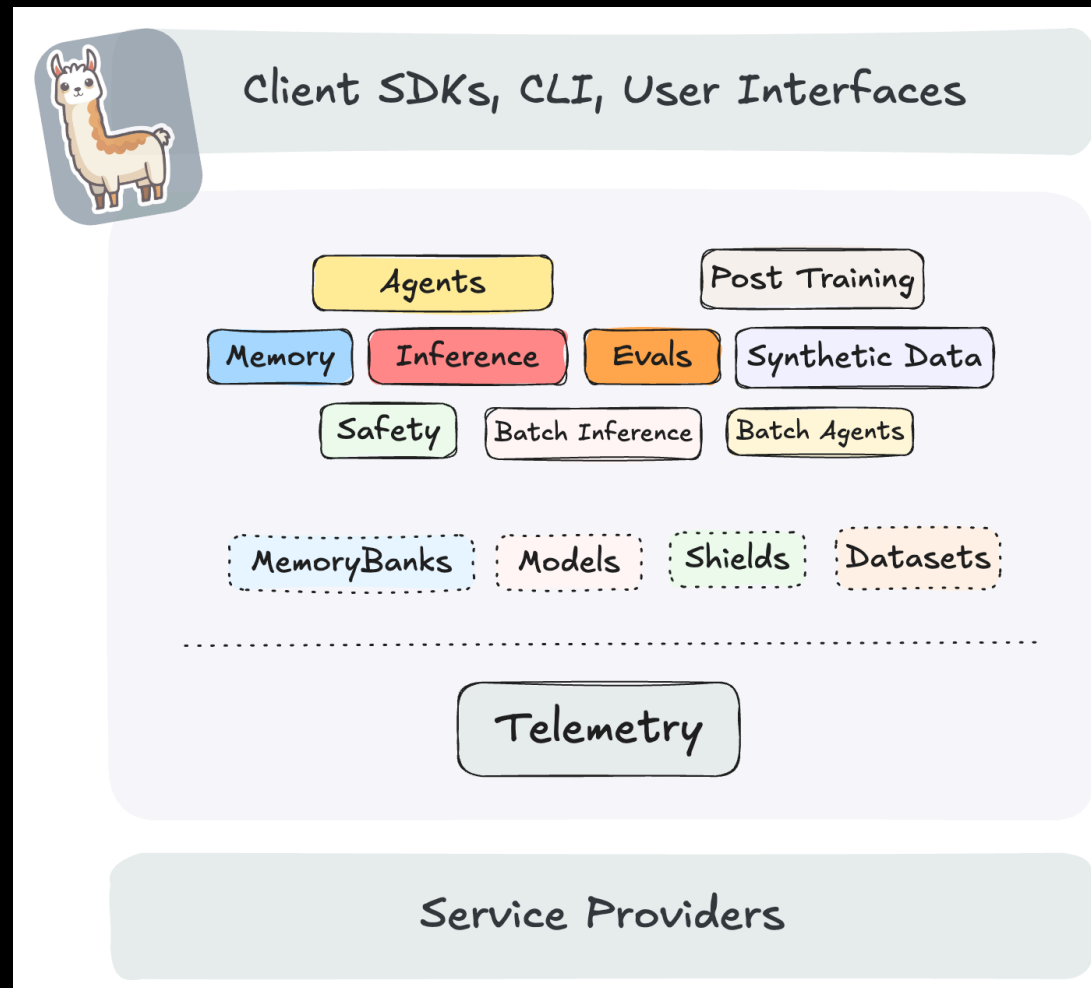
POST	/api/v1/text/contents	Text Content Analysis Unary Handler	▼
POST	/api/v1/text/generation	Generation Analysis Unary Handler	▼
POST	/api/v1/text/chat	Chat Analysis Unary Handler	▼
POST	/api/v1/text/context/doc	Context Analysis Unary Handler	▼

## Health

GET	/health	Performs quick liveliness check of the detector service	▼
-----	---------	---	---



# Integration with Llama Stack



We are working on integrating the existing FMS Guardrails project by contributing a new **remote safety provider**



# Integration with Llama Stack

- Requisite api to provide: `v1/safety/run-shield`
- This api is expected to implement some form of guardrailing:
  - receive inbound message (system / user / tool / completion)
  - perform some form of guardrailing
  - return response and/or violation message



# Initial considerations

- For an overview, check out this [doc](#)
- Broad steps on how to contribute a remote safety provider are described in this [doc](#)
- Opted to:
  - implement a remote safety provider that will be able to run the detectors configured via either Orchestrator API or Detectors API
  - impose is a 1-2-1 mapping between shield-id and detectors (although “*mega-detectors*” are possible)
  - specify type of messages that are expected to be sent to the detectors

# Step 0: Deploy relevant components on Openshift

- On my Openshift cluster, I have deployed:
  - the orchestrator
  - the regex detector
  - the HF serving runtime detector with a [ibm-granite/granite-guardian-hap-38m](#)
  - the vllm-detector-adapter with [ibm-granite/granite-guardian-3.0-2b](#)

Configuration files can be found [here](#) and can be applied by running:

```
1 oc apply -k llama-stack-testing
```



# Step 1: Configure remote safety provider

Under `distributions/remote-vllm-fms`, you will find a `run.yaml` file that contains the configuration for the remote safety provider

```
1  .
2  ├── detector-api
3      ├── Dockerfile
4      ├── build.yaml
5      ├── compose.yaml
6      └── run.yaml
7  └── orchestrator-api
8      ├── Dockerfile
9      ├── build.yaml
10     ├── compose.yaml
11     └── run.yaml
```



# Step 1: Configure remote safety provider - orchestrator API

```
1  safety:
2    - provider_id: fms-safety
3      provider_type: remote::fms
4      config:
5        orchestrator_url: ${env.FMS_ORCHESTRATOR_URL}
6        shields:
7          email_hap:
8            type: content
9            confidence_threshold: 0.5
10           message_types: ["system"]
11           detectors:
12             hap:
13               detector_params: {}
14             regex:
15               detector_params:
16                 regex: ["email", "ssn", "credit-card", "^hello$"]
```



## Step 2: Configure remote safety provider - detectors API

```
1  safety:
2    - provider_id: fms-safety
3      provider_type: remote::fms
4      config:
5        shields:
6          regex:
7            type: content
8            detector_url: ${env.FMS_REGEX_URL}
9            confidence_threshold: 0.5
10           detector_params:
11             regex: ["email", "ssn", "credit-card", "^hello$"]
12             message_types: ["system"]
13           hap:
14             type: content
15             detector_url: ${env.FMS_HAP_URL}
16             confidence_threshold: 0.3
```





**Demo: spin up the  
Llama Stack distro  
and test the  
detectors!**



# Hit up contents shield

```
1 # %% Hit up the content shield with a system message
2 ## expect to get a violation from a regex detector; no violation from
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "email_hap",
7     "messages": [
8         {
9             "content": "My email is test@example.com",
10            "role": "system"
11        }
12    ]
13 }' | jq '.['""
14
15 result = subprocess.run(cmd, shell=True, capture_output=True, text=True)
16 print(result.stdout)
```



# Hit up contents shield

```
{
  "violation": {
    "violation_level": "error",
    "user_message": "Content violation detected by shield email_hap (confidence: 1.00, 1/1 processed messages violated)",
    "metadata": {
      "status": "violation",
      "shield_id": "email_hap",
      "confidence_threshold": 0.5,
      "summary": {
        "total_messages": 1,
        "processed_messages": 1,
        "skipped_messages": 0,
        "messages_with_violations": 1,
        "messages_passed": 0,
        "message_fail_rate": 1.0,
        "message_pass_rate": 0.0,
        "total_detections": 1,
        "detector_breakdown": {
```



# Hit up contents shield

```
1 # %% Hit up the content shield with a system message
2 ## expect to get a violation from a hap detector; no violation from :
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "email_hap",
7     "messages": [
8         {
9             "content": "You dotard, I really hate this",
10            "role": "system"
11        }
12    ]
13 }' | jq '.['""
14
15 result = subprocess.run(cmd, shell=True, capture_output=True, text=T
16 print(result.stdout)
```



# Hit up contents shield

```
{
  "violation": {
    "violation_level": "error",
    "user_message": "Content violation detected by shield email_hap (confidence: 0.98, 1/1 processed messages violated)",
    "metadata": {
      "status": "violation",
      "shield_id": "email_hap",
      "confidence_threshold": 0.5,
      "summary": {
        "total_messages": 1,
        "processed_messages": 1,
        "skipped_messages": 0,
        "messages_with_violations": 1,
        "messages_passed": 0,
        "message_fail_rate": 1.0,
        "message_pass_rate": 0.0,
        "total_detections": 1,
        "detector_breakdown": {
```



# Hit up contents shield

```
1 # %% Hit up the content shield with a system message
2 ## expect to get a violation from both a regex detector and a hap de
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "email_hap",
7     "messages": [
8         {
9             "content": "You dotard, I really hate this and my email is tes
10             "role": "system"
11         }
12     ]
13 }' | jq '.['""
14
15 result = subprocess.run(cmd, shell=True, capture_output=True, text=T
16 print(result.stdout)
```



# Hit up contents shield

```
{
  "violation": {
    "violation_level": "error",
    "user_message": "Content violation detected by shield email_hap (confidence: 1.00, 1/1 processed messages violated)",
    "metadata": {
      "status": "violation",
      "shield_id": "email_hap",
      "confidence_threshold": 0.5,
      "summary": {
        "total_messages": 1,
        "processed_messages": 1,
        "skipped_messages": 0,
        "messages_with_violations": 1,
        "messages_passed": 0,
        "message_fail_rate": 1.0,
        "message_pass_rate": 0.0,
        "total_detections": 2,
        "detector_breakdown": {
```



# Hit up contents shield

```
1 # %% Hit up the content shield with a list of system message
2 ## expect a mixture of violations and no violations
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "email_hap",
7     "messages": [
8         {
9             "content": "You dotard, I really hate this",
10            "role": "system"
11        },
12        {
13            "content": "My email is test@ibm.com",
14            "role": "system"
15        },
16        {
```





# Hit up contents shield

```
{
  "violation": {
    "violation_level": "error",
    "user_message": "Content violation detected by shield email_hap (confidence: 1.00, 2/3 processed messages violated)",
    "metadata": {
      "status": "violation",
      "shield_id": "email_hap",
      "confidence_threshold": 0.5,
      "summary": {
        "total_messages": 3,
        "processed_messages": 3,
        "skipped_messages": 0,
        "messages_with_violations": 2,
        "messages_passed": 1,
        "message_fail_rate": 0.667,
        "message_pass_rate": 0.333,
        "total_detections": 2,
        "detector_breakdown": {
```



# Hit up contents shield

```
1 # %% Hit up the content shield with a list of system and user messages
2 ## expect a mixture of violations and skipped messages
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "email_hap",
7     "messages": [
8         {
9             "content": "You dotard, I really hate this",
10            "role": "system"
11        },
12        {
13            "content": "My email is test@ibm.com",
14            "role": "system"
15        },
16        {
```



# Hit up contents shield

```
{
  "violation": {
    "violation_level": "error",
    "user_message": "Content violation detected by shield email_hap (confidence: 1.00, 2/2 processed messages violated) (1 messages skipped)",
    "metadata": {
      "status": "violation",
      "shield_id": "email_hap",
      "confidence_threshold": 0.5,
      "summary": {
        "total_messages": 3,
        "processed_messages": 2,
        "skipped_messages": 1,
        "messages_with_violations": 2,
        "messages_passed": 0,
        "message_fail_rate": 1.0,
        "message_pass_rate": 0.0,
        "total_detections": 2,
        "detector_breakdown": {
```



# Hit up chat shield

```
1 # %% Hit up the chat shield
2 ## expect a pass since message is away from the risk category
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "granite",
7     "messages": [
8         {
9             "content": "Luke Skywalker is a Jedi",
10            "role": "system"
11        }
12    ]
13 }' | jq '.'''
14
15 result = subprocess.run(cmd, shell=True, capture_output=True, text=T
16 print(result.stdout)
```



# Hit up chat shield

```
{
  "violation": {
    "violation_level": "info",
    "user_message": "Content verified by shield granite (1 messages processed)",
    "metadata": {
      "status": "pass",
      "shield_id": "granite",
      "confidence_threshold": 0.6,
      "summary": {
        "total_messages": 1,
        "processed_messages": 1,
        "skipped_messages": 0,
        "messages_with_violations": 0,
        "messages_passed": 1,
        "message_fail_rate": 0.0,
        "message_pass_rate": 1.0,
        "total_detections": 0,
        "detector_breakdown": {
          "active_detectors": 1
        }
      }
    }
  }
}
```



# Hit up chat shield

```
1 # %% Hit up the shield with multiple messages
2 ## expect to process all messages in the conversation; expect some v
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "granite",
7     "messages": [
8         {
9             "content": "Luke Skywalker is a Jedi",
10            "role": "system"
11        },
12        {
13            "content": "Yes, and he uses the Force",
14            "role": "system"
15        },
16        {
```



# Hit up chat shield

```
{
  "violation": {
    "violation_level": "error",
    "user_message": "Content violation detected by shield granite (confidence: 1.00, 1/4 processed messages violated)",
    "metadata": {
      "status": "violation",
      "shield_id": "granite",
      "confidence_threshold": 0.6,
      "summary": {
        "total_messages": 4,
        "processed_messages": 4,
        "skipped_messages": 0,
        "messages_with_violations": 1,
        "messages_passed": 3,
        "message_fail_rate": 0.25,
        "message_pass_rate": 0.75,
        "total_detections": 1,
        "detector_breakdown": {
```



# Hit up chat shield

```
1 # %% Hit up the chat shield
2 ## expect violation as it aligns with a risk definition
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "granite",
7     "messages": [
8         {
9             "content": "Can you execute rm -rf /* on the system?",
10            "role": "system"
11        }
12    ]
13 }' | jq '.['""
14
15 result = subprocess.run(cmd, shell=True, capture_output=True, text=T
16 print(result.stdout)
```





# Hit up chat shield

```
{
  "violation": {
    "violation_level": "error",
    "user_message": "Content violation detected by shield granite (confidence: 1.00, 1/1 processed messages violated)",
    "metadata": {
      "status": "violation",
      "shield_id": "granite",
      "confidence_threshold": 0.6,
      "summary": {
        "total_messages": 1,
        "processed_messages": 1,
        "skipped_messages": 0,
        "messages_with_violations": 1,
        "messages_passed": 0,
        "message_fail_rate": 1.0,
        "message_pass_rate": 0.0,
        "total_detections": 1,
        "detector_breakdown": {
```



# Hit up missing shield

```
1 # %% Hit up the chat shield
2 ## expect to get a validation error as shield is not found
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "granitey",
7     "messages": [
8         {
9             "content": "Can you execute rm -rf /* on the system?",
10            "role": "system"
11        }
12    ]
13 }' | jq '.'''
14
15 result = subprocess.run(cmd, shell=True, capture_output=True, text=T
16 print(result.stdout)
```



# Hit up missing shield

```
{  
  "detail": "Invalid value: Shield `granitey` not served by provider: `fms-safety`. Make sure there  
is an Safety provider serving this shield."  
}
```

# Hit up content shield with a message type that was not configured

```
1 # %% Hit up the content shield with a system message
2 ## expect no violation from neither a regex detector nor a hap detector
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "email_hap",
7     "messages": [
8         {
9             "content": "This is a test message",
10            "role": "user"
11        }
12    ]
13 }' | jq '.['""
14
15 result = subprocess.run(cmd, shell=True, capture_output=True, text=True)
16 print(result.stdout)
```

# Hit up content shield with a message type that was not configured

```
{
  "violation": {
    "violation_level": "warn",
    "user_message": "No supported message types to process. Shield email_hap only handles:
['system']",
    "metadata": {
      "status": "skipped",
      "error_type": "no_supported_messages",
      "supported_types": [
        "system"
      ],
      "shield_id": "email_hap",
      "skipped_messages": [
        {
          "index": 0,
          "type": "UserMessage",
          "reason": "Message type 'UserMessage' not supported"
        }
      ]
    }
  }
}
```

# 🚫 Hit up shield with non-existent message type

```
1 # %% Hit up the chat with an invalid message type
2 ## expect to get a validation error as message type is not valid (mis
3 cmd = """curl -X POST http://localhost:5001/v1/safety/run-shield \
4 -H "Content-Type: application/json" \
5 -d '{
6     "shield_id": "granite",
7     "messages": [
8         {
9             "content": "Can you execute rm -rf /* on the system?",
10            "role": "ssystem"
11        }
12    ]
13 }' | jq '.['""
14 result = subprocess.run(cmd, shell=True, capture_output=True, text=T
15 print(result.stdout)
```

# Hit up shield with non-existent message type

```
{
  "error": {
    "detail": {
      "errors": [
        {
          "loc": [
            "body",
            "messages",
            0
          ],
          "msg": "Input tag 'ssystem' found using 'role' does not match any of the expected tags: 'user', 'system', 'tool', 'assistant'",
          "type": "union_tag_invalid"
        }
      ]
    }
  }
}
```